

# Predicting the Sale Price of Bulldozers using Machine Learning



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [6]: pd , np ,plt , sns
```

```
Out[6]: (<module 'pandas' from 'D:\\Projects\\bulldozer price prediction\\env\\Lib\\site-packages\\pandas\\__init__.py'>,
<module 'numpy' from 'D:\\Projects\\bulldozer price prediction\\env\\Lib\\site-packages\\numpy\\__init__.py'>,
<module 'matplotlib.pyplot' from 'D:\\Projects\\bulldozer price prediction\\env\\Lib\\site-packages\\matplotlib\\pyplot.py'>,
<module 'seaborn' from 'D:\\Projects\\bulldozer price prediction\\env\\Lib\\site-packages\\seaborn\\__init__.py'>)
```

```
In [7]: df = pd.read_csv("data/TrainAndValid.csv" , low_memory = False)
```

```
In [8]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 412698 entries, 0 to 412697

Data columns (total 53 columns):

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	saledate	412698 non-null	object
10	fiModelDesc	412698 non-null	object
11	fiBaseModel	412698 non-null	object
12	fiSecondaryDesc	271971 non-null	object
13	fiModelSeries	58667 non-null	object
14	fiModelDescriptor	74816 non-null	object
15	ProductSize	196093 non-null	object
16	fiProductClassDesc	412698 non-null	object
17	state	412698 non-null	object
18	ProductGroup	412698 non-null	object
19	ProductGroupDesc	412698 non-null	object
20	Drive_System	107087 non-null	object
21	Enclosure	412364 non-null	object
22	Forks	197715 non-null	object
23	Pad_Type	81096 non-null	object
24	Ride_Control	152728 non-null	object
25	Stick	81096 non-null	object
26	Transmission	188007 non-null	object
27	Turbocharged	81096 non-null	object
28	Blade_Extension	25983 non-null	object
29	Blade_Width	25983 non-null	object
30	Enclosure_Type	25983 non-null	object
31	Engine_Horsepower	25983 non-null	object
32	Hydraulics	330133 non-null	object
33	Pushblock	25983 non-null	object
34	Ripper	106945 non-null	object
35	Scarifier	25994 non-null	object
36	Tip_Control	25983 non-null	object
37	Tire_Size	97638 non-null	object
38	Coupler	220679 non-null	object
39	Coupler_System	44974 non-null	object
40	Grouser_Tracks	44875 non-null	object
41	Hydraulics_Flow	44875 non-null	object
42	Track_Type	102193 non-null	object
43	Undercarriage_Pad_Width	102916 non-null	object
44	Stick_Length	102261 non-null	object
45	Thumb	102332 non-null	object
46	Pattern_Changer	102261 non-null	object
47	Grouser_Type	102193 non-null	object
48	Backhoe_Mounting	80712 non-null	object
49	Blade_Type	81875 non-null	object
50	Travel_Controls	81877 non-null	object


```
51 Differential_Type      71564 non-null  object
52 Steering_Controls      71522 non-null  object
dtypes: float64(3), int64(5), object(45)
memory usage: 166.9+ MB
```

```
In [9]: df.isna().sum()
```

```
Out[9]: SalesID          0
        SalePrice        0
        MachineID        0
        ModelID          0
        datasource       0
        auctioneerID     20136
        YearMade         0
        MachineHoursCurrentMeter 265194
        UsageBand        339028
        saledate         0
        fiModelDesc      0
        fiBaseModel      0
        fiSecondaryDesc   140727
        fiModelSeries     354031
        fiModelDescriptor 337882
        ProductSize       216605
        fiProductClassDesc 0
        state            0
        ProductGroup      0
        ProductGroupDesc  0
        Drive_System      305611
        Enclosure         334
        Forks            214983
        Pad_Type          331602
        Ride_Control      259970
        Stick             331602
        Transmission      224691
        Turbocharged      331602
        Blade_Extension   386715
        Blade_Width       386715
        Enclosure_Type    386715
        Engine_Horsepower 386715
        Hydraulics        82565
        Pushblock         386715
        Ripper            305753
        Scarifier         386704
        Tip_Control       386715
        Tire_Size         315060
        Coupler           192019
        Coupler_System    367724
        Grouser_Tracks    367823
        Hydraulics_Flow   367823
        Track_Type        310505
        Undercarriage_Pad_Width 309782
        Stick_Length      310437
        Thumb            310366
        Pattern_Changer   310437
        Grouser_Type      310505
        Backhoe_Mounting  331986
        Blade_Type        330823
        Travel_Controls   330821
        Differential_Type  341134
        Steering_Controls 341176
        dtype: int64
```

```
In [10]: fig , ax = plt.subplots()
ax.scatter(df["saledate"][:1000] , df["SalePrice"][:1000])
```

Out[10]: <matplotlib.collections.PathCollection at 0x25f8ed3e3c0>




```
In [11]: fig, ax = plt.subplots(figsize=(10,6))

ax.scatter(df["saledate"][:1000], df["SalePrice"][:1000], alpha=0.5)

ax.set_xlabel("Sale Date")
ax.set_ylabel("Sale Price")
ax.set_title("Bulldozer Sale Price over Time (first 1000 samples)")

plt.xticks(rotation=45) # rotate dates for readability
plt.show()
```




```
In [12]: df.saledate
```

```
Out[12]: 0      11/16/2006 0:00
1      3/26/2004 0:00
2      2/26/2004 0:00
3      5/19/2011 0:00
4      7/23/2009 0:00
...
412693   3/7/2012 0:00
412694   1/28/2012 0:00
412695   1/28/2012 0:00
412696   3/7/2012 0:00
412697   1/28/2012 0:00
Name: saledate, Length: 412698, dtype: object
```

```
In [13]: # Plot histogram
df["SalePrice"].hist() # bins = number of intervals
```

```
Out[13]: <Axes: >
```



## Prasing Data

```
In [16]: df = pd.read_csv("data/TrainAndValid.csv" , low_memory= False , parse_dates=
```

```
In [38]: df
```

```
Out[38]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Y
<b>0</b>	1139246	66000.0	999089	3157	121	3.0	
<b>1</b>	1139248	57000.0	117657	77	121	3.0	
<b>2</b>	1139249	10000.0	434808	7009	121	3.0	
<b>3</b>	1139251	38500.0	1026470	332	121	3.0	
<b>4</b>	1139253	11000.0	1057373	17311	121	3.0	
...	...	...	...	...	...	...	...
<b>412693</b>	6333344	10000.0	1919201	21435	149	2.0	
<b>412694</b>	6333345	10500.0	1882122	21436	149	2.0	
<b>412695</b>	6333347	12500.0	1944213	21435	149	2.0	
<b>412696</b>	6333348	10000.0	1794518	21435	149	2.0	
<b>412697</b>	6333349	13000.0	1944743	21436	149	2.0	

412698 rows × 53 columns

```
In [17]: df.saledate[:1000]
```

```
Out[17]:
```

0	2006-11-16
1	2004-03-26
2	2004-02-26
3	2011-05-19
4	2009-07-23
...	...
995	2009-07-16
996	2007-06-14
997	2005-09-22
998	2005-07-28
999	2011-06-16

Name: saledate, Length: 1000, dtype: datetime64[ns]

```
In [44]: df["saledate"]
```



```
Out[44]: 0      2006-11-16
1      2004-03-26
2      2004-02-26
3      2011-05-19
4      2009-07-23
...
412693 2012-03-07
412694 2012-01-28
412695 2012-01-28
412696 2012-03-07
412697 2012-01-28
Name: saledate, Length: 412698, dtype: datetime64[ns]
```

```
In [18]: fig, ax = plt.subplots()
ax.scatter(df["saledate"][:1000], df["SalePrice"][:1000])
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x25f922191d0>
```




```
In [19]: fig, ax = plt.subplots(figsize=(10,6))

ax.scatter(df["saledate"][:1000], df["SalePrice"][:1000], alpha=0.5)

ax.set_xlabel("Sale Date")
ax.set_ylabel("Sale Price")
ax.set_title("Bulldozer Sale Price over Time (first 1000 samples)")

plt.xticks(rotation=45) # rotate dates for readability
plt.show()
```



```
In [20]: df.head().T
```

Out[20]:

	0	1	2	3
<b>SalesID</b>	1139246	1139248	1139249	1139251
<b>SalePrice</b>	66000.0	57000.0	10000.0	38500.0
<b>MachineID</b>	999089	117657	434808	1026470
<b>ModelID</b>	3157	77	7009	332
<b>datasource</b>	121	121	121	121
<b>auctioneerID</b>	3.0	3.0	3.0	3.0
<b>YearMade</b>	2004	1996	2001	2001
<b>MachineHoursCurrentMeter</b>	68.0	4640.0	2838.0	3486.0
<b>UsageBand</b>	Low	Low	High	High
<b>saledate</b>	2006-11-16 00:00:00	2004-03-26 00:00:00	2004-02-26 00:00:00	2011-05-19 00:00:00
<b>fiModelDesc</b>	521D	950FII	226	PC120-6E
<b>fiBaseModel</b>	521	950	226	PC120
<b>fiSecondaryDesc</b>	D	F	NaN	NaN
<b>fiModelSeries</b>	NaN	II	NaN	-6E
<b>fiModelDescriptor</b>	NaN	NaN	NaN	NaN
<b>ProductSize</b>	NaN	Medium	NaN	Small
<b>fiProductClassDesc</b>	Wheel Loader - 110.0 to 120.0 Horsepower	Wheel Loader - 150.0 to 175.0 Horsepower	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...
<b>state</b>	Alabama	North Carolina	New York	Texas
<b>ProductGroup</b>	WL	WL	SSL	TEX
<b>ProductGroupDesc</b>	Wheel Loader	Wheel Loader	Skid Steer Loaders	Track Excavators
<b>Drive_System</b>	NaN	NaN	NaN	NaN
<b>Enclosure</b>	EROPS w AC	EROPS w AC	OROPS	EROPS w AC
<b>Forks</b>	None or Unspecified	None or Unspecified	None or Unspecified	NaN
<b>Pad_Type</b>	NaN	NaN	NaN	NaN
<b>Ride_Control</b>	None or Unspecified	None or Unspecified	NaN	NaN
<b>Stick</b>	NaN	NaN	NaN	NaN
<b>Transmission</b>	NaN	NaN	NaN	NaN

	0	1	2	3
<b>Turbocharged</b>	NaN	NaN	NaN	NaN
<b>Blade_Extension</b>	NaN	NaN	NaN	NaN
<b>Blade_Width</b>	NaN	NaN	NaN	NaN
<b>Enclosure_Type</b>	NaN	NaN	NaN	NaN
<b>Engine_Horsepower</b>	NaN	NaN	NaN	NaN
<b>Hydraulics</b>	2 Valve	2 Valve	Auxiliary	2 Valve
<b>Pushblock</b>	NaN	NaN	NaN	NaN
<b>Ripper</b>	NaN	NaN	NaN	NaN
<b>Scarifier</b>	NaN	NaN	NaN	NaN
<b>Tip_Control</b>	NaN	NaN	NaN	NaN
<b>Tire_Size</b>	None or Unspecified	23.5	NaN	NaN
<b>Coupler</b>	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified
<b>Coupler_System</b>	NaN	NaN	None or Unspecified	NaN
<b>Grouser_Tracks</b>	NaN	NaN	None or Unspecified	NaN
<b>Hydraulics_Flow</b>	NaN	NaN	Standard	NaN
<b>Track_Type</b>	NaN	NaN	NaN	NaN
<b>Undercarriage_Pad_Width</b>	NaN	NaN	NaN	NaN
<b>Stick_Length</b>	NaN	NaN	NaN	NaN
<b>Thumb</b>	NaN	NaN	NaN	NaN
<b>Pattern_Changer</b>	NaN	NaN	NaN	NaN
<b>Grouser_Type</b>	NaN	NaN	NaN	NaN
<b>Backhoe_Mounting</b>	NaN	NaN	NaN	NaN
<b>Blade_Type</b>	NaN	NaN	NaN	NaN
<b>Travel_Controls</b>	NaN	NaN	NaN	NaN
<b>Differential_Type</b>	Standard	Standard	NaN	NaN
<b>Steering_Controls</b>	Conventional	Conventional	NaN	NaN

## Sort Dataframe by Saledate

```
In [23]: df.sort_values(by=["saledate"], inplace = True , ascending = True)
```

```
In [24]: df["saledate"].head(50)
```

```
Out[24]: 205615    1989-01-17
233186    1989-01-31
142491    1989-01-31
115536    1989-01-31
92301     1989-01-31
115892    1989-01-31
134080    1989-01-31
92294     1989-01-31
31494     1989-01-31
140922    1989-01-31
66337     1989-01-31
92531     1989-01-31
82122     1989-01-31
92256     1989-01-31
145670    1989-01-31
92780     1989-01-31
238373    1989-01-31
127132    1989-01-31
115102    1989-01-31
32317     1989-01-31
238656    1989-01-31
52508     1989-01-31
127923    1989-01-31
127521    1989-01-31
152689    1989-01-31
82165     1989-01-31
78445     1989-01-31
62665     1989-01-31
113454    1989-01-31
113547    1989-01-31
28820     1989-01-31
168619    1989-01-31
115957    1989-01-31
205782    1989-01-31
114830    1989-01-31
127735    1989-01-31
78382     1989-01-31
127674    1989-01-31
28603     1989-01-31
78278     1989-01-31
231507    1989-01-31
169757    1989-01-31
92803     1989-01-31
75832     1989-01-31
88803     1989-01-31
75378     1989-01-31
169297    1989-01-31
280078    1989-01-31
140257    1989-01-31
128751    1989-01-31
Name: saledate, dtype: datetime64[ns]
```

## Make a copy of original Data frame

```
In [25]: df_tmp = df.copy()
```

```
In [59]: df_tmp
```

```
Out[59]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Y
<b>205615</b>	1646770	9500.0	1126363	8434	132	18.0	
<b>274835</b>	1821514	14000.0	1194089	10150	132	99.0	
<b>141296</b>	1505138	50000.0	1473654	4139	132	99.0	
<b>212552</b>	1671174	16000.0	1327630	8591	132	99.0	
<b>62755</b>	1329056	22000.0	1336053	4089	132	99.0	
...	...	...	...	...	...	...	
<b>410879</b>	6302984	16000.0	1915521	5266	149	99.0	
<b>412476</b>	6324811	6000.0	1919104	19330	149	99.0	
<b>411927</b>	6313029	16000.0	1918416	17244	149	99.0	
<b>407124</b>	6266251	55000.0	509560	3357	149	99.0	
<b>409203</b>	6283635	34000.0	1869284	4701	149	99.0	

412698 rows × 53 columns

## add Datetime parameter for "SaleDate" column

```
In [26]: df_tmp["saleyear"] = df_tmp.saledate.dt.year  
df_tmp["saleMonth"] = df_tmp.saledate.dt.month  
df_tmp["saleDay"] = df_tmp.saledate.dt.day  
df_tmp["saleDayofWeek"] = df_tmp.saledate.dt.dayofweek  
df_tmp["saleDayofyear"] = df_tmp.saledate.dt.dayofyear
```

```
In [7]: df_tmp.head().T
```

Out[7]:

	205615	274835	141296	212552
<b>SalesID</b>	1646770	1821514	1505138	1671174
<b>SalePrice</b>	9500.0	14000.0	50000.0	16000.0
<b>MachineID</b>	1126363	1194089	1473654	1327630
<b>ModelID</b>	8434	10150	4139	8591
<b>datasource</b>	132	132	132	132
<b>auctioneerID</b>	18.0	99.0	99.0	99.0
<b>YearMade</b>	1974	1980	1978	1980
<b>MachineHoursCurrentMeter</b>	NaN	NaN	NaN	NaN
<b>UsageBand</b>	NaN	NaN	NaN	NaN
<b>saledate</b>	1989-01-17 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00
<b>fiModelDesc</b>	TD20	A66	D7G	A62
<b>fiBaseModel</b>	TD20	A66	D7	A62
<b>fiSecondaryDesc</b>	NaN	NaN	G	NaN
<b>fiModelSeries</b>	NaN	NaN	NaN	NaN
<b>fiModelDescriptor</b>	NaN	NaN	NaN	NaN
<b>ProductSize</b>	Medium	NaN	Large	NaN
<b>fiProductClassDesc</b>	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	Wheel Loader - Unidentified
<b>state</b>	Texas	Florida	Florida	Florida
<b>ProductGroup</b>	TTT	WL	TTT	WL
<b>ProductGroupDesc</b>	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader
<b>Drive_System</b>	NaN	NaN	NaN	NaN
<b>Enclosure</b>	OROPS	OROPS	OROPS	EROPS
<b>Forks</b>	NaN	None or Unspecified	NaN	None or Unspecified
<b>Pad_Type</b>	NaN	NaN	NaN	NaN
<b>Ride_Control</b>	NaN	None or Unspecified	NaN	None or Unspecified
<b>Stick</b>	NaN	NaN	NaN	NaN
<b>Transmission</b>	Direct Drive	NaN	Standard	NaN

	205615	274835	141296	212552
<b>Turbocharged</b>	NaN	NaN	NaN	NaN
<b>Blade_Extension</b>	NaN	NaN	NaN	NaN
<b>Blade_Width</b>	NaN	NaN	NaN	NaN
<b>Enclosure_Type</b>	NaN	NaN	NaN	NaN
<b>Engine_Horsepower</b>	NaN	NaN	NaN	NaN
<b>Hydraulics</b>	2 Valve	2 Valve	2 Valve	2 Valve
<b>Pushblock</b>	NaN	NaN	NaN	NaN
<b>Ripper</b>	None or Unspecified	NaN	None or Unspecified	NaN
<b>Scarifier</b>	NaN	NaN	NaN	NaN
<b>Tip_Control</b>	NaN	NaN	NaN	NaN
<b>Tire_Size</b>	NaN	None or Unspecified	NaN	None or Unspecified
<b>Coupler</b>	NaN	None or Unspecified	NaN	None or Unspecified
<b>Coupler_System</b>	NaN	NaN	NaN	NaN
<b>Grouser_Tracks</b>	NaN	NaN	NaN	NaN
<b>Hydraulics_Flow</b>	NaN	NaN	NaN	NaN
<b>Track_Type</b>	NaN	NaN	NaN	NaN
<b>Undercarriage_Pad_Width</b>	NaN	NaN	NaN	NaN
<b>Stick_Length</b>	NaN	NaN	NaN	NaN
<b>Thumb</b>	NaN	NaN	NaN	NaN
<b>Pattern_Changer</b>	NaN	NaN	NaN	NaN
<b>Grouser_Type</b>	NaN	NaN	NaN	NaN
<b>Backhoe_Mounting</b>	None or Unspecified	NaN	None or Unspecified	NaN
<b>Blade_Type</b>	Straight	NaN	Straight	NaN
<b>Travel_Controls</b>	None or Unspecified	NaN	None or Unspecified	NaN
<b>Differential_Type</b>	NaN	Standard	NaN	Standard
<b>Steering_Controls</b>	NaN	Conventional	NaN	Conventional
<b>saleyear</b>	1989	1989	1989	1989
<b>saleMonth</b>	1	1	1	1
<b>saleDay</b>	17	31	31	31
<b>saleDayofWeek</b>	1	1	1	1



	205615	274835	141296	212552
saleDayofyear	17	31	31	31

## Now we going to remove sale date

```
In [27]: df_tmp.drop("saledate" , axis = 1 , inplace = True)
```

```
In [28]: df_tmp.columns
```

```
Out[28]: Index(['SalesID', 'SalePrice', 'MachineID', 'ModelID', 'datasource',
               'auctioneerID', 'YearMade', 'MachineHoursCurrentMeter', 'UsageBand',
               'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc', 'fiModelSeries',
               'fiModelDescriptor', 'ProductSize', 'fiProductClassDesc', 'state',
               'ProductGroup', 'ProductGroupDesc', 'Drive_System', 'Enclosure',
               'Forks', 'Pad_Type', 'Ride_Control', 'Stick', 'Transmission',
               'Turbocharged', 'Blade_Extension', 'Blade_Width', 'Enclosure_Type',
               'Engine_Horsepower', 'Hydraulics', 'Pushblock', 'Ripper', 'Scarifie
r',
               'Tip_Control', 'Tire_Size', 'Coupler', 'Coupler_System',
               'Grouser_Tracks', 'Hydraulics_Flow', 'Track_Type',
               'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb', 'Pattern_Change
r',
               'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type', 'Travel_Controls',
               'Differential_Type', 'Steering_Controls', 'saleyear', 'saleMonth',
               'saleDay', 'saleDayofWeek', 'saleDayofyear'],
              dtype='object')
```

```
In [29]: # Check the values of different columns
df_tmp.state.value_counts()
```

```
Out[29]: state
Florida      67320
Texas        53110
California    29761
Washington    16222
Georgia       14633
Maryland      13322
Mississippi   13240
Ohio          12369
Illinois      11540
Colorado      11529
New Jersey    11156
North Carolina 10636
Tennessee     10298
Alabama       10292
Pennsylvania  10234
South Carolina 9951
Arizona       9364
New York      8639
Connecticut   8276
Minnesota     7885
Missouri      7178
Nevada        6932
Louisiana     6627
Kentucky      5351
Maine         5096
Indiana       4124
Arkansas      3933
New Mexico    3631
Utah          3046
Unspecified   2801
Wisconsin     2745
New Hampshire 2738
Virginia      2353
Idaho         2025
Oregon        1911
Michigan      1831
Wyoming       1672
Montana       1336
Iowa          1336
Oklahoma      1326
Nebraska      866
West Virginia 840
Kansas        667
Delaware      510
North Dakota  480
Alaska        430
Massachusetts 347
Vermont       300
South Dakota  244
Hawaii        118
Rhode Island   83
Puerto Rico   42
Washington DC  2
Name: count, dtype: int64
```

## 5. Modelling

```
In [30]: from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor(n_jobs = -1 , random_state=42)
```

```
In [15]: df
```

```
Out[15]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Y
<b>205615</b>	1646770	9500.0	1126363	8434	132	18.0	
<b>274835</b>	1821514	14000.0	1194089	10150	132	99.0	
<b>141296</b>	1505138	50000.0	1473654	4139	132	99.0	
<b>212552</b>	1671174	16000.0	1327630	8591	132	99.0	
<b>62755</b>	1329056	22000.0	1336053	4089	132	99.0	
...	...	...	...	...	...	...	
<b>410879</b>	6302984	16000.0	1915521	5266	149	99.0	
<b>412476</b>	6324811	6000.0	1919104	19330	149	99.0	
<b>411927</b>	6313029	16000.0	1918416	17244	149	99.0	
<b>407124</b>	6266251	55000.0	509560	3357	149	99.0	
<b>409203</b>	6283635	34000.0	1869284	4701	149	99.0	

412698 rows × 53 columns

```
In [31]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 412698 entries, 205615 to 409203

Data columns (total 53 columns):

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	saledate	412698 non-null	datetime64[ns]
10	fiModelDesc	412698 non-null	object
11	fiBaseModel	412698 non-null	object
12	fiSecondaryDesc	271971 non-null	object
13	fiModelSeries	58667 non-null	object
14	fiModelDescriptor	74816 non-null	object
15	ProductSize	196093 non-null	object
16	fiProductClassDesc	412698 non-null	object
17	state	412698 non-null	object
18	ProductGroup	412698 non-null	object
19	ProductGroupDesc	412698 non-null	object
20	Drive_System	107087 non-null	object
21	Enclosure	412364 non-null	object
22	Forks	197715 non-null	object
23	Pad_Type	81096 non-null	object
24	Ride_Control	152728 non-null	object
25	Stick	81096 non-null	object
26	Transmission	188007 non-null	object
27	Turbocharged	81096 non-null	object
28	Blade_Extension	25983 non-null	object
29	Blade_Width	25983 non-null	object
30	Enclosure_Type	25983 non-null	object
31	Engine_Horsepower	25983 non-null	object
32	Hydraulics	330133 non-null	object
33	Pushblock	25983 non-null	object
34	Ripper	106945 non-null	object
35	Scarifier	25994 non-null	object
36	Tip_Control	25983 non-null	object
37	Tire_Size	97638 non-null	object
38	Coupler	220679 non-null	object
39	Coupler_System	44974 non-null	object
40	Grouser_Tracks	44875 non-null	object
41	Hydraulics_Flow	44875 non-null	object
42	Track_Type	102193 non-null	object
43	Undercarriage_Pad_Width	102916 non-null	object
44	Stick_Length	102261 non-null	object
45	Thumb	102332 non-null	object
46	Pattern_Changer	102261 non-null	object
47	Grouser_Type	102193 non-null	object
48	Backhoe_Mounting	80712 non-null	object
49	Blade_Type	81875 non-null	object
50	Travel_Controls	81877 non-null	object

```
51 Differential_Type          71564 non-null  object
52 Steering_Controls          71522 non-null  object
dtypes: datetime64[ns](1), float64(3), int64(5), object(44)
memory usage: 170.0+ MB
```

```
In [32]: df.isna().sum()
```

```

Out[32]: SalesID          0
         SalePrice        0
         MachineID        0
         ModelID          0
         datasource        0
         auctioneerID      20136
         YearMade          0
         MachineHoursCurrentMeter 265194
         UsageBand        339028
         saledate          0
         fiModelDesc       0
         fiBaseModel       0
         fiSecondaryDesc    140727
         fiModelSeries      354031
         fiModelDescriptor  337882
         ProductSize        216605
         fiProductClassDesc 0
         state              0
         ProductGroup       0
         ProductGroupDesc   0
         Drive_System       305611
         Enclosure          334
         Forks              214983
         Pad_Type           331602
         Ride_Control       259970
         Stick              331602
         Transmission       224691
         Turbocharged       331602
         Blade_Extension    386715
         Blade_Width        386715
         Enclosure_Type     386715
         Engine_Horsepower  386715
         Hydraulics         82565
         Pushblock          386715
         Ripper             305753
         Scarifier          386704
         Tip_Control        386715
         Tire_Size          315060
         Coupler            192019
         Coupler_System     367724
         Grouser_Tracks     367823
         Hydraulics_Flow    367823
         Track_Type         310505
         Undercarriage_Pad_Width 309782
         Stick_Length       310437
         Thumb              310366
         Pattern_Changer    310437
         Grouser_Type       310505
         Backhoe_Mounting   331986
         Blade_Type         330823
         Travel_Controls    330821
         Differential_Type   341134
         Steering_Controls   341176
         dtype: int64

```

## Convert String to categories

```
In [27]: df_tmp.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 412698 entries, 205615 to 409203

Data columns (total 57 columns):

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	fiModelDesc	412698 non-null	object
10	fiBaseModel	412698 non-null	object
11	fiSecondaryDesc	271971 non-null	object
12	fiModelSeries	58667 non-null	object
13	fiModelDescriptor	74816 non-null	object
14	ProductSize	196093 non-null	object
15	fiProductClassDesc	412698 non-null	object
16	state	412698 non-null	object
17	ProductGroup	412698 non-null	object
18	ProductGroupDesc	412698 non-null	object
19	Drive_System	107087 non-null	object
20	Enclosure	412364 non-null	object
21	Forks	197715 non-null	object
22	Pad_Type	81096 non-null	object
23	Ride_Control	152728 non-null	object
24	Stick	81096 non-null	object
25	Transmission	188007 non-null	object
26	Turbocharged	81096 non-null	object
27	Blade_Extension	25983 non-null	object
28	Blade_Width	25983 non-null	object
29	Enclosure_Type	25983 non-null	object
30	Engine_Horsepower	25983 non-null	object
31	Hydraulics	330133 non-null	object
32	Pushblock	25983 non-null	object
33	Ripper	106945 non-null	object
34	Scarifier	25994 non-null	object
35	Tip_Control	25983 non-null	object
36	Tire_Size	97638 non-null	object
37	Coupler	220679 non-null	object
38	Coupler_System	44974 non-null	object
39	Grouser_Tracks	44875 non-null	object
40	Hydraulics_Flow	44875 non-null	object
41	Track_Type	102193 non-null	object
42	Undercarriage_Pad_Width	102916 non-null	object
43	Stick_Length	102261 non-null	object
44	Thumb	102332 non-null	object
45	Pattern_Changer	102261 non-null	object
46	Grouser_Type	102193 non-null	object
47	Backhoe_Mounting	80712 non-null	object
48	Blade_Type	81875 non-null	object
49	Travel_Controls	81877 non-null	object
50	Differential_Type	71564 non-null	object



```
51 Steering_Controls          71522 non-null  object
52 saleyear                   412698 non-null  int32
53 saleMonth                   412698 non-null  int32
54 saleDay                     412698 non-null  int32
55 saleDayofWeek               412698 non-null  int32
56 saleDayofyear               412698 non-null  int32
dtypes: float64(3), int32(5), int64(5), object(44)
memory usage: 174.7+ MB
```

```
In [33]: pd.api.types.is_string_dtype(df_tmp["UsageBand"])
```

```
Out[33]: False
```

```
In [34]: for label , content in df_tmp.items():
          if pd.api.types.is_string_dtype(content):

              print(label)
          #if pd.api.types.is_string_dtype(df_tmp())
```

```
fiModelDesc
fiBaseModel
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
```

```
In [35]: for label ,content in df_tmp.items():
          if pd.api.types.is_object_dtype(content):
              df_tmp[label] = content.astype("category").cat.as_ordered()
```

```
In [36]: df_tmp.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 412698 entries, 205615 to 409203

Data columns (total 57 columns):

#	Column	Non-Null Count	Dtype
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	category
9	fiModelDesc	412698 non-null	category
10	fiBaseModel	412698 non-null	category
11	fiSecondaryDesc	271971 non-null	category
12	fiModelSeries	58667 non-null	category
13	fiModelDescriptor	74816 non-null	category
14	ProductSize	196093 non-null	category
15	fiProductClassDesc	412698 non-null	category
16	state	412698 non-null	category
17	ProductGroup	412698 non-null	category
18	ProductGroupDesc	412698 non-null	category
19	Drive_System	107087 non-null	category
20	Enclosure	412364 non-null	category
21	Forks	197715 non-null	category
22	Pad_Type	81096 non-null	category
23	Ride_Control	152728 non-null	category
24	Stick	81096 non-null	category
25	Transmission	188007 non-null	category
26	Turbocharged	81096 non-null	category
27	Blade_Extension	25983 non-null	category
28	Blade_Width	25983 non-null	category
29	Enclosure_Type	25983 non-null	category
30	Engine_Horsepower	25983 non-null	category
31	Hydraulics	330133 non-null	category
32	Pushblock	25983 non-null	category
33	Ripper	106945 non-null	category
34	Scarifier	25994 non-null	category
35	Tip_Control	25983 non-null	category
36	Tire_Size	97638 non-null	category
37	Coupler	220679 non-null	category
38	Coupler_System	44974 non-null	category
39	Grouser_Tracks	44875 non-null	category
40	Hydraulics_Flow	44875 non-null	category
41	Track_Type	102193 non-null	category
42	Undercarriage_Pad_Width	102916 non-null	category
43	Stick_Length	102261 non-null	category
44	Thumb	102332 non-null	category
45	Pattern_Changer	102261 non-null	category
46	Grouser_Type	102193 non-null	category
47	Backhoe_Mounting	80712 non-null	category
48	Blade_Type	81875 non-null	category
49	Travel_Controls	81877 non-null	category
50	Differential_Type	71564 non-null	category

```

51 Steering_Controls          71522 non-null    category
52 saleyear                   412698 non-null   int32
53 saleMonth                  412698 non-null   int32
54 saleDay                    412698 non-null   int32
55 saleDayofWeek              412698 non-null   int32
56 saleDayofyear              412698 non-null   int32
dtypes: category(44), float64(3), int32(5), int64(5)
memory usage: 55.4 MB

```

```
In [37]: df_tmp.state.cat.codes
```

```

Out[37]: 205615    43
          233186    8
          142491    8
          115536    8
           92301    8
          ..
          409901    4
          405777    4
          411889    4
          411890    4
          409203    4
Length: 412698, dtype: int8

```

## Thanks for categories

Fill Missing values

Fill numerical missing values first

```
In [38]: for label, content in df_tmp.items():
          if pd.api.types.is_numeric_dtype(content):
              print(df_tmp[label].isna().sum(), label)
```

```

0 SalesID
0 SalePrice
0 MachineID
0 ModelID
0 datasource
20136 auctioneerID
0 YearMade
265194 MachineHoursCurrentMeter
0 saleyear
0 saleMonth
0 saleDay
0 saleDayofWeek
0 saleDayofyear

```

```
In [61]: df_tmp.ModelID
```

```
Out[61]: 205615      8434
         274835     10150
         141296      4139
         212552      8591
         62755      4089
         ...
         410879      5266
         412476     19330
         411927     17244
         407124      3357
         409203      4701
Name: ModelID, Length: 412698, dtype: int64
```

```
In [39]: # Check for which numeric columns have null values
```

```
for label , content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

auctioneerID  
MachineHoursCurrentMeter

```
In [40]: # Fill numeric rows with the median
```

```
for label , content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

auctioneerID  
MachineHoursCurrentMeter

```
In [41]: # Fill numeric rows with the median
```

```
for label , content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
            df_tmp[label + "_is_missing"] = pd.isnull(content)
            df_tmp[label] = content.fillna(content.median())
```

auctioneerID  
MachineHoursCurrentMeter

```
In [42]: # Fill numeric rows with the median
```

```
for label , content in df_tmp.items():  
    if pd.api.types.is_numeric_dtype(content):  
        if pd.isnull(content).sum():  
            print(label)
```

```
In [43]: df_tmp.auctioneerID_is_missing.value_counts()
```

```
Out[43]: auctioneerID_is_missing  
False    392562  
True      20136  
Name: count, dtype: int64
```

## Filling and turning categorical variables into numbers

```
In [44]: for label ,content in df_tmp.items():  
        if not pd.api.types.is_numeric_dtype(content):  
            print(label)
```

UsageBand  
fiModelDesc  
fiBaseModel  
fiSecondaryDesc  
fiModelSeries  
fiModelDescriptor  
ProductSize  
fiProductClassDesc  
state  
ProductGroup  
ProductGroupDesc  
Drive\_System  
Enclosure  
Forks  
Pad\_Type  
Ride\_Control  
Stick  
Transmission  
Turbocharged  
Blade\_Extension  
Blade\_Width  
Enclosure\_Type  
Engine\_Horsepower  
Hydraulics  
Pushblock  
Ripper  
Scarifier  
Tip\_Control  
Tire\_Size  
Coupler  
Coupler\_System  
Grouser\_Tracks  
Hydraulics\_Flow  
Track\_Type  
Undercarriage\_Pad\_Width  
Stick\_Length  
Thumb  
Pattern\_Changer  
Grouser\_Type  
Backhoe\_Mounting  
Blade\_Type  
Travel\_Controls  
Differential\_Type  
Steering\_Controls

```
In [45]: # Turn categorical variables into numbers
```

```
for label , content in df_tmp.items():  
    if not pd.api.types.is_numeric_dtype(content):  
        # Add binary column to indicate  
        df_tmp[label+"_is_missing"] = pd.isnull(content)  
  
        df_tmp[label] = pd.Categorical(content).codes+1
```

```
In [46]: df_tmp.isna().sum()
```

```
Out[46]: SalesID          0
         SalePrice       0
         MachineID       0
         ModelID         0
         datasource      0
         ..
         Backhoe_Mounting_is_missing  0
         Blade_Type_is_missing       0
         Travel_Controls_is_missing  0
         Differential_Type_is_missing 0
         Steering_Controls_is_missing 0
         Length: 103, dtype: int64
```

```
In [47]: df.isna().sum()
```

```

Out[47]: SalesID          0
         SalePrice        0
         MachineID        0
         ModelID          0
         datasource        0
         auctioneerID     20136
         YearMade          0
         MachineHoursCurrentMeter 265194
         UsageBand        339028
         saledate          0
         fiModelDesc       0
         fiBaseModel       0
         fiSecondaryDesc   140727
         fiModelSeries     354031
         fiModelDescriptor 337882
         ProductSize       216605
         fiProductClassDesc 0
         state             0
         ProductGroup      0
         ProductGroupDesc  0
         Drive_System      305611
         Enclosure         334
         Forks             214983
         Pad_Type          331602
         Ride_Control      259970
         Stick             331602
         Transmission      224691
         Turbocharged      331602
         Blade_Extension   386715
         Blade_Width       386715
         Enclosure_Type    386715
         Engine_Horsepower 386715
         Hydraulics        82565
         Pushblock         386715
         Ripper            305753
         Scarifier         386704
         Tip_Control       386715
         Tire_Size         315060
         Coupler           192019
         Coupler_System    367724
         Grouser_Tracks    367823
         Hydraulics_Flow   367823
         Track_Type        310505
         Undercarriage_Pad_Width 309782
         Stick_Length      310437
         Thumb             310366
         Pattern_Changer   310437
         Grouser_Type      310505
         Backhoe_Mounting  331986
         Blade_Type        330823
         Travel_Controls   330821
         Differential_Type  341134
         Steering_Controls  341176
         dtype: int64

```



```
In [48]: df_tmp.head()
```

```
Out[48]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	Y
<b>205615</b>	1646770	9500.0	1126363	8434	132	18.0	
<b>274835</b>	1821514	14000.0	1194089	10150	132	99.0	
<b>141296</b>	1505138	50000.0	1473654	4139	132	99.0	
<b>212552</b>	1671174	16000.0	1327630	8591	132	99.0	
<b>62755</b>	1329056	22000.0	1336053	4089	132	99.0	

5 rows × 103 columns

```
In [51]: %%time
```

```
model = RandomForestRegressor(n_jobs = -1 , random_state=42)
model.fit(df_tmp.drop("SalePrice" , axis =1) ,df_tmp["SalePrice"])
```

CPU times: total: 31min 4s

Wall time: 3min 5s

```
Out[51]:
```

RandomForestRegressor ⓘ ?

▶ Parameters

```
In [52]: %%time
```

```
model.score(df_tmp.drop("SalePrice" , axis =1) ,df_tmp["SalePrice"])
```

CPU times: total: 38.8 s

Wall time: 4.1 s

```
Out[52]: 0.9875966080326709
```

## Splitting data into train/validation sets

```
In [53]: df_val = df_tmp[df_tmp.saleyear == 2012]
df_train = df_tmp[df_tmp.saleyear != 2012]
```

```
In [55]: len(df_train) , len(df_val)
```

```
Out[55]: (401125, 11573)
```

```
In [56]: len(df_train) + len(df_val)
```

```
Out[56]: 412698
```

```
In [57]: x_train , y_train = df_train.drop("SalePrice" , axis = 1) , df_train.SalePri
```

```
In [58]: x_valid , y_valid = df_val.drop("SalePrice" , axis = 1) , df_val.SalePrice
```

```
x_train.shape , y_train.shape , x_valid.shape , y_valid.shape
```

```
Out[58]: ((401125, 102), (401125,), (11573, 102), (11573,))
```

## Building and evaluation Funtion

```
In [59]: # Create evaluation function
```

```
from sklearn.metrics import mean_squared_log_error , mean_absolute_error
from sklearn.metrics import r2_score
```

```
def rmsle(y_test,y_preds):
```

```
    """
```

```
    Calculate root mean squard log error
```

```
    """
```

```
    return np.sqrt(mean_squared_log_error(y_test , y_preds))
```

```
# Create function to evaluate model on a few different levels
```

```
def show_scores(model):
```

```
    train_preds = model.predict(x_train)
```

```
    val_preds = model.predict(x_valid)
```

```
    scores = {"Training MAE" : mean_absolute_error(y_train, train_preds),
              "Valid MAE" : mean_absolute_error(y_valid , val_preds),
              "Training RMSLE" : rmsle(y_train , train_preds),
              "Valid RMSLE" : rmsle(y_valid , val_preds) ,
              "Training R2" : r2_score(y_train , train_preds),
              "Valid R2" : r2_score(y_valid , val_preds)
              }
```

```
    return scores
```

## Testing our model on subset

```
%%time
```

```
model = RandomForestRegressor(n_jobs=-1 , random_state=42)
```

```
model.fit(x_train , y_train)
```

```
In [60]: # CHnage max sample value
```

```
model = RandomForestRegressor(n_jobs=-1 , random_state=42 , max_samples = 10000)
```

```
model
```

Out[60]:

▼ RandomForestRegressor ⓘ ?  
► Parameters

In [61]: `x_train.shape[0]`

Out[61]: 401125

In [62]: `%%time`

```
model.fit(x_train , y_train)
```

CPU times: total: 59.1 s

Wall time: 6.07 s

Out[62]:

▼ RandomForestRegressor ⓘ ?  
► Parameters

In [63]: `show_scores(model)`

```
Out[63]: {'Training MAE': 5548.995840324088,  
          'Valid MAE': 7179.6961392897265,  
          'Training RMSLE': np.float64(0.25737726780537257),  
          'Valid RMSLE': np.float64(0.29404344200903443),  
          'Training R2': 0.8610738743845617,  
          'Valid R2': 0.8320179198265637}
```

In [64]: `%%time`

```
from sklearn.model_selection import RandomizedSearchCV  
  
rf_grid = {"n_estimators":
```

CPU times: total: 0 ns

Wall time: 31.2 µs

## Train a model with best hyperparameters

In [65]: `%%time`

```
ideal_model = RandomForestRegressor(n_estimators=40 , min_samples_leaf=1, mi  
  
ideal_model.fit(x_train , y_train)
```

CPU times: total: 5min 19s

Wall time: 32.5 s

Out[65]:

▼ RandomForestRegressor ⓘ ?  
► Parameters

```
In [92]: show_scores(ideal_model)
```

```
Out[92]: {'Training MAE': 2954.3860545002144,
          'Valid MAE': 5934.273509420097,
          'Training RMSLE': np.float64(0.14445948886340432),
          'Valid RMSLE': np.float64(0.24607663981354877),
          'Training R2': 0.9588911792610285,
          'Valid R2': 0.8824941147924968}
```

```
In [66]: # Scores for ideal model(trained on all data)
show_scores(ideal_model)
```

```
Out[66]: {'Training MAE': 2956.4887000110048,
          'Valid MAE': 5957.457704159022,
          'Training RMSLE': np.float64(0.14448478225577735),
          'Valid RMSLE': np.float64(0.24544663828024385),
          'Training R2': 0.9587936671021189,
          'Valid R2': 0.8820814857765893}
```

## Make predictions on test data

```
In [67]: # Import the test data
```

```
test_df = pd.read_csv("Data/Test.csv" , low_memory=False , parse_dates=["sale"
```

```
In [68]: len(test_df.columns)
```

```
Out[68]: 52
```

## Make predictions

```
''' test_preds = ideal_model.predict(test_df) '''
```

This has error because this is not same as the training set Now we are going to preprocess the data and make sure its same as ...

```
In [69]: def preprocess_data(test_df):
        """
        Perform transformation on df and returns transformed df.
        """
        test_df["saleyear"]      = test_df.saledate.dt.year
        test_df["saleMonth"]     = test_df.saledate.dt.month
        test_df["saleDay"]       = test_df.saledate.dt.day
        test_df["saleDayofWeek"] = test_df.saledate.dt.dayofweek
        test_df["saleDayofyear"] = test_df.saledate.dt.dayofyear

        test_df.drop("saledate" , axis =1 , inplace=True)'''

        # Fill the numeric rows with median
        for label , content in test_df.items():
            if pd.api.types.is_numeric_dtype(content):
```

```
    if pd.isnull(content).sum():
        print(label)
        test_df[label + "_is_missing"] = pd.isnull(content)
        test_df[label] = content.fillna(content.median())

#Filled categorical missing data and turned in numbers

    if not pd.api.types.is_numeric_dtype(content):
        # Add binary column to indicate
        test_df[label + "_is_missing"] = pd.isnull(content)

        test_df[label] = pd.Categorical(content).codes + 1

    return test_df
```

```
In [70]: test_df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 12457 entries, 0 to 12456

Data columns (total 52 columns):

#	Column	Non-Null Count	Dtype
0	SalesID	12457 non-null	int64
1	MachineID	12457 non-null	int64
2	ModelID	12457 non-null	int64
3	datasource	12457 non-null	int64
4	auctioneerID	12457 non-null	int64
5	YearMade	12457 non-null	int64
6	MachineHoursCurrentMeter	2129 non-null	float64
7	UsageBand	1834 non-null	object
8	saledate	12457 non-null	datetime64[ns]
9	fiModelDesc	12457 non-null	object
10	fiBaseModel	12457 non-null	object
11	fiSecondaryDesc	8482 non-null	object
12	fiModelSeries	2006 non-null	object
13	fiModelDescriptor	3024 non-null	object
14	ProductSize	6048 non-null	object
15	fiProductClassDesc	12457 non-null	object
16	state	12457 non-null	object
17	ProductGroup	12457 non-null	object
18	ProductGroupDesc	12457 non-null	object
19	Drive_System	2759 non-null	object
20	Enclosure	12455 non-null	object
21	Forks	6308 non-null	object
22	Pad_Type	2108 non-null	object
23	Ride_Control	4241 non-null	object
24	Stick	2108 non-null	object
25	Transmission	4818 non-null	object
26	Turbocharged	2108 non-null	object
27	Blade_Extension	651 non-null	object
28	Blade_Width	651 non-null	object
29	Enclosure_Type	651 non-null	object
30	Engine_Horsepower	651 non-null	object
31	Hydraulics	10315 non-null	object
32	Pushblock	651 non-null	object
33	Ripper	2704 non-null	object
34	Scarifier	651 non-null	object
35	Tip_Control	651 non-null	object
36	Tire_Size	2778 non-null	object
37	Coupler	7601 non-null	object
38	Coupler_System	2066 non-null	object
39	Grouser_Tracks	2066 non-null	object
40	Hydraulics_Flow	2066 non-null	object
41	Track_Type	3394 non-null	object
42	Undercarriage_Pad_Width	3398 non-null	object
43	Stick_Length	3394 non-null	object
44	Thumb	3395 non-null	object
45	Pattern_Changer	3394 non-null	object
46	Grouser_Type	3394 non-null	object
47	Backhoe_Mounting	2051 non-null	object
48	Blade_Type	2058 non-null	object
49	Travel_Controls	2058 non-null	object
50	Differential_Type	2129 non-null	object

51 Steering\_Controls 2129 non-null object  
dtypes: datetime64[ns](1), float64(1), int64(6), object(44)  
memory usage: 4.9+ MB

In [71]: preprocess\_data(test\_df)

MachineHoursCurrentMeter

Out[71]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	M
0	1227829	1006309	3168	121	3	1999	
1	1227844	1022817	7271	121	3	1000	
2	1227847	1031560	22805	121	3	2004	
3	1227848	56204	1269	121	3	2006	
4	1227863	1053887	22312	121	3	2005	
...	...	...	...	...	...	...	...
12452	6643171	2558317	21450	149	2	2008	
12453	6643173	2558332	21434	149	2	2005	
12454	6643184	2558342	21437	149	2	1000	
12455	6643186	2558343	21437	149	2	2006	
12456	6643196	2558346	21446	149	2	2008	

12457 rows × 98 columns

In [85]: test\_df

Out[85]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	M
0	1227829	1006309	3168	121	False	1999	
1	1227844	1022817	7271	121	False	1000	
2	1227847	1031560	22805	121	False	2004	
3	1227848	56204	1269	121	False	2006	
4	1227863	1053887	22312	121	False	2005	
...	...	...	...	...	...	...	...
12452	6643171	2558317	21450	149	False	2008	
12453	6643173	2558332	21434	149	False	2005	
12454	6643184	2558342	21437	149	False	1000	
12455	6643186	2558343	21437	149	False	2006	
12456	6643196	2558346	21446	149	False	2008	

12457 rows × 102 columns

```
In [72]: test_df.isna().sum()
```

```
Out[72]: SalesID          0
MachineID        0
ModelID          0
datasource       0
auctioneerID     0
..
Backhoe_Mounting_is_missing  0
Blade_Type_is_missing      0
Travel_Controls_is_missing  0
Differential_Type_is_missing 0
Steering_Controls_is_missing 0
Length: 98, dtype: int64
```

```
In [79]: test_df.head()
```

```
Out[79]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	Machir
0	1227829	1006309	3168	121	False	1999	
1	1227844	1022817	7271	121	False	1000	
2	1227847	1031560	22805	121	False	2004	
3	1227848	56204	1269	121	False	2006	
4	1227863	1053887	22312	121	False	2005	

5 rows × 102 columns

```
In [58]: x_train.head()
```

```
Out[58]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	I
205615	1646770	1126363	8434	132	18.0	1974	
274835	1821514	1194089	10150	132	99.0	1980	
141296	1505138	1473654	4139	132	99.0	1978	
212552	1671174	1327630	8591	132	99.0	1980	
62755	1329056	1336053	4089	132	99.0	1984	

5 rows × 102 columns

```
In [74]: # GOING TO FIND THE COLUMN DIFFER
set(x_train.columns) - set(test_df.columns)
```



```
Out[74]: {'auctioneerID_is_missing',
          'saleDay',
          'saleDayofWeek',
          'saleDayofyear',
          'saleMonth',
          'saleyear'}
```

```
In [75]: ## adjust df_test to have auctioneer id

test_df["auctioneerID_is_missing"] = False
```

```
In [76]: len(x_valid) , len(test_df)
```

```
Out[76]: (11573, 12457)
```

```
In [76]: test_df
```

```
Out[76]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	M
0	1227829	1006309	3168	121	False	1999	
1	1227844	1022817	7271	121	False	1000	
2	1227847	1031560	22805	121	False	2004	
3	1227848	56204	1269	121	False	2006	
4	1227863	1053887	22312	121	False	2005	
...	...	...	...	...	...	...	...
12452	6643171	2558317	21450	149	False	2008	
12453	6643173	2558332	21434	149	False	2005	
12454	6643184	2558342	21437	149	False	1000	
12455	6643186	2558343	21437	149	False	2006	
12456	6643196	2558346	21446	149	False	2008	

12457 rows × 102 columns

## Make predictions on test data

```
test_preds = ideal_model.predict(test_df)
```

```
In [77]: test_df = preprocess_data(test_df)
```

```
In [90]: test_df
```

Out[90]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	M
<b>0</b>	1227829	1006309	3168	121	False	1999	
<b>1</b>	1227844	1022817	7271	121	False	1000	
<b>2</b>	1227847	1031560	22805	121	False	2004	
<b>3</b>	1227848	56204	1269	121	False	2006	
<b>4</b>	1227863	1053887	22312	121	False	2005	
<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	<b>...</b>	
<b>12452</b>	6643171	2558317	21450	149	False	2008	
<b>12453</b>	6643173	2558332	21434	149	False	2005	
<b>12454</b>	6643184	2558342	21437	149	False	1000	
<b>12455</b>	6643186	2558343	21437	149	False	2006	
<b>12456</b>	6643196	2558346	21446	149	False	2008	

12457 rows × 102 columns

```
In [78]: # Preprocess test data
test_df_proc = preprocess_data(test_df)

# Match the columns to training data
test_df_proc = test_df_proc.reindex(columns=x_train.columns, fill_value=0)

# Now predict
test_preds = ideal_model.predict(test_df_proc)
```

```
In [79]: len(test_preds)
```

Out[79]: 12457

```
In [99]: test_df_proc
```

Out[99]:

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	M
0	1227829	1006309	3168	121	False	1999	
1	1227844	1022817	7271	121	False	1000	
2	1227847	1031560	22805	121	False	2004	
3	1227848	56204	1269	121	False	2006	
4	1227863	1053887	22312	121	False	2005	
...	...	...	...	...	...	...	...
12452	6643171	2558317	21450	149	False	2008	
12453	6643173	2558332	21434	149	False	2005	
12454	6643184	2558342	21437	149	False	1000	
12455	6643186	2558343	21437	149	False	2006	
12456	6643196	2558346	21446	149	False	2008	

12457 rows × 102 columns

In [80]: *# Format predictions into the same format*

```
df_preds = pd.DataFrame()

df_preds["SalesID"] = test_df["SalesID"]
df_preds["SalesPrice"] = test_preds
```

In [105... df\_preds

Out[105...

	SalesID	SalesPrice
0	1227829	17623.337125
1	1227844	14566.296570
2	1227847	46662.254410
3	1227848	71305.266295
4	1227863	61762.999424
...	...	...
12452	6643171	40469.885910
12453	6643173	12196.277617
12454	6643184	11964.850733
12455	6643186	16342.165338
12456	6643196	27119.990440

12457 rows × 2 columns

```
In [81]: df_preds.to_csv("test_prediction.csv" , index = False)
```


```
In [82]: len(ideal_model.feature_importances_)
```

```
Out[82]: 102
```

```
In [83]: # Helper function for plotting feature importance
def plot_features(columns, importances, n=20):
    # Create a DataFrame with features and their importance
    df = (
        pd.DataFrame({
            "feature": columns,
            "importance": importances
        })
        .sort_values("importance", ascending=False)
        .reset_index(drop=True)
    )

    # Plot
    fig, ax = plt.subplots(figsize=(10, 6))
    ax.barh(df["feature"][:n][::-1], df["importance"][:n][::-1]) # horizontal
    ax.set_xlabel("Feature Importance")
    ax.set_title(f"Top {n} Important Features")
    plt.show()
```

```
In [118... plot_features(x_train.columns , ideal_model.feature_importances_)
```



```
In [119... x_train.head()
```

Out[119...

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	I
<b>205615</b>	1646770	1126363	8434	132	18.0	1974	
<b>274835</b>	1821514	1194089	10150	132	99.0	1980	
<b>141296</b>	1505138	1473654	4139	132	99.0	1978	
<b>212552</b>	1671174	1327630	8591	132	99.0	1980	
<b>62755</b>	1329056	1336053	4089	132	99.0	1984	

5 rows × 102 columns

In [ ]: